

Specialist Diploma in Artificial Intelligence

AI Applications with Deep Learning

Creating Your Own Chatbot

Sub topics

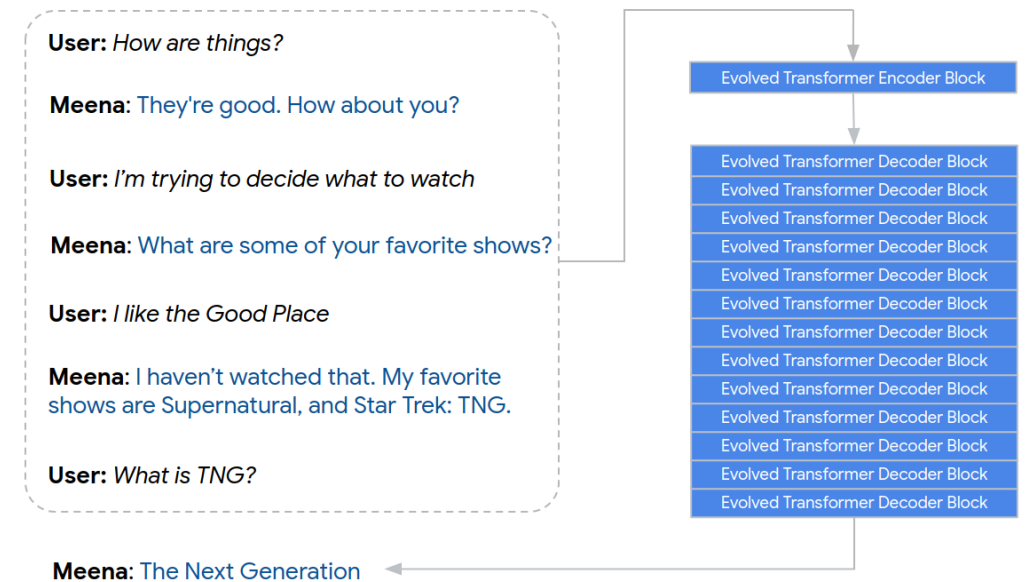
- **What is a Chatbot?**
- **Transformers with the HuggingFace Library**
- **Intent Recognition**
- **Custom Named Entity Recognition**
- **Putting It Altogether**
- **Other Advances in Natural Language Processing**

Creating Your Own Chatbot

WHAT IS A CHATBOT?

Open Domain Chatbot

- Multi-Turned Open Domain Chatbots
 - Meena Chatbot by Google
 - Uses seq2seq architecture
 - Has 2.6B parameters.
 - Trained on 341GB dataset of social media conversations



<https://ai.googleblog.com/2020/01/towards-conversational-agent-that-can.html>

Google's Meena Chatbot



<https://www.youtube.com/watch?v=cA0vHYzGR8s>

A Simple Chatbot – Closed Domain

- What we will try to build:
 - Closed Domain = Chatbot meets specific business objectives
- For each text said by the user:
 - Use AI to Classify Intent (Intent Classification)
 - Use AI to Extract Entities (Entity Recognition)
 - Write Rules to Process Intent and Entities
 - Write Code to Display Reply to User

Creating Your Own Chatbot

TRANSFORMERS WITH HUGGINGFACE LIBRARY

Recall...

- The HuggingFace library:
 - A popular Deep Learning Natural Language Processing library
 - Includes many variations of the Transformer architecture
 - Includes many pipelines for various Natural Language Processing tasks
 - Can run on-premise (non-cloud)





Supported NLP Tasks

- Sequence Classification

Topic modelling, text classification, positive/negative sentiment analysis, intent recognition

“This is a great movie”





POSITIVE

Supported NLP Tasks

- Extractive Question & Answer

Answers a question by extracting an answer from FAQs, customer reviews, social media posts, etc.

“How do I contact the sales team?”



“You can call 011-2233-4444”





The blue text is generated as a reply by the transformer. This is likely extracted from a paragraph of text in its training data

Supported NLP Tasks


- Text Generation

Generates a human-like text that keeps the writing as coherent as possible

“As far as I am concerned,
I will”



“As far as I am concerned, I will
**be the first to admit that I am a
not fan of the idea of...**”





The blue text is generated by
the transformer

Supported NLP Tasks

- Named Entity Recognition

Classify each word or token in a sentence to identify what parts of speech, or what properties they belong to.

“Hugging Face Inc. is a
company based in New
York City”



“Hugging Face Inc.”
“New York City.”

ORGANIZATION
LOCATION





The transformer decides
which parts of the text belong
to certain entities

Supported NLP Tasks

- Summarization

Summarizes a paragraph of text.

New York (CNN) When Liana Barrientos was 23 years old, she got married in Westchester County, New York. ... A year later, she got married again in Westchester County, but to a different man and without divorcing her first husband. ... Only 18 days after that marriage, she got hitched yet again. Then, Barrientos declared "I do" five more times, sometimes only within two weeks of each other. ... In 2010, she married once more, this time in the Bronx. In an application for a marriage license, she stated it was her "first and only" marriage. ... Barrientos, now 39, is facing two criminal counts of "offering a false instrument for filing in the first degree," referring to her false statements on the ... 2010 marriage license application, according to court documents. ... Prosecutors said the marriages were part of an immigration scam. ... On Friday, she pleaded not guilty at State Supreme Court in the Bronx, according to her attorney, Christopher Wright, who declined to comment further. ... After leaving court, Barrientos was arrested and charged with theft of service and criminal trespass for allegedly sneaking into the New York subway through an emergency exit, said Detective ... Annette Markowski, a police spokeswoman. In total, Barrientos has been married 10 times, with nine of her marriages occurring between 1999 and 2002. ... All occurred either in Westchester County, Long Island, New Jersey or the Bronx. She is believed to still be married to four men, and at one time, she was married to eight men at once, prosecutors say. ... Prosecutors said the immigration scam involved some of her husbands, who filed for permanent residence status shortly after the marriages. ... Any divorces happened only after such filings were approved. It was unclear whether any of the men will be prosecuted. ... The case was referred to the Bronx District Attorney's Office by Immigration and Customs Enforcement and the Department of Homeland Security's ... Investigation Division. Seven of the men are from so-called "red-flagged" countries, including Egypt, Turkey, Georgia, Pakistan and Mali. ... Her eighth husband, Rashid Rajput, was deported in 2006 to his native Pakistan after an investigation by the Joint Terrorism Task Force. ... If convicted, Barrientos faces up to four years in prison. Her next court appearance is scheduled for May 18.





Liana Barrientos, 39, is charged with two counts of "offering a false instrument for filing in the first degree" In total, she has been married 10 times, with nine of her marriages occurring between 1999 and 2002. She is believed to still be married to four men.

Supported NLP Tasks

- Translation

Takes a text in one language and translates to another language.

“Hugging Face is a
technology company based
in New York and Paris”



“Hugging Face ist ein
Technologieunternehmen mit
Sitz in New York und Paris.”

Using a 3rd-Party Library

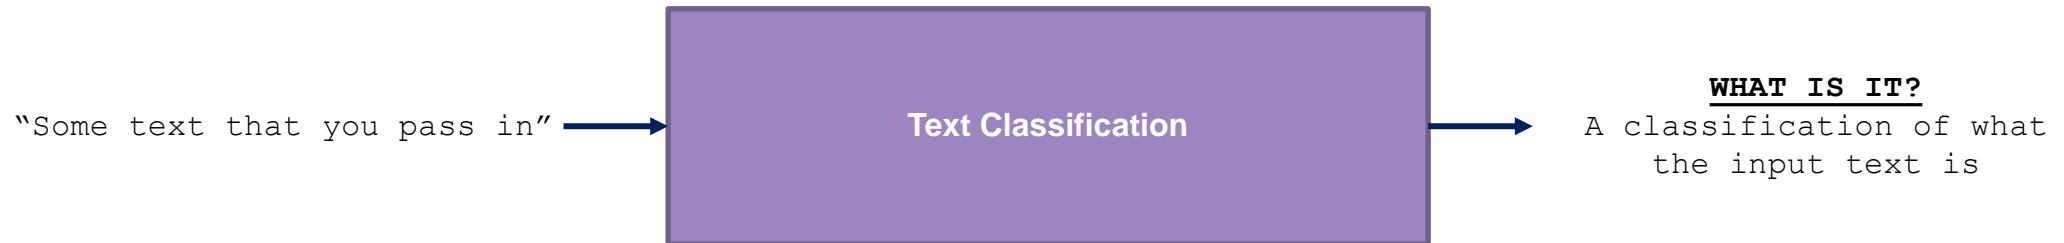
- Using third party framework / tools are common when applying Deep Learning to your problem statement
- Importantly, we need to know how to use that framework to:
 - Construct our training / test data to the model for training
 - Construct an input to the model for prediction / inferencing

Creating Your Own Chatbot

INTENT RECOGNITION – ZERO SHOT CLASSIFICATION

Text Classification


- Intent recognition can be achieved with “Sequence / Text Classification”.
- What we are actually doing:



Some Use Cases



Social Media Monitoring
(Complaints, Threats)



Automated Tagging of
User Content
(Tag Suggestions, Categorization)



Support Ticket Classification /
Intent Recognition



Document Classification
and Search



Sentiment Analysis
(Branding, Public Relations)

Dataset

- We want to build a chatbot that interacts with customers on **flight enquiries** and other flight-related matters.
- Assuming that we use a combination of the following datasets:
 - <https://www.kaggle.com/hassanamin/atis-airlinetravelinformationsystem>
ATIS that captured customers' requests and classified each of them
 - <https://github.com/alyssaong1/botframework-smalltalk/blob/master/smalltalkkb.tsv>
A small talk chatbot dataset with greetings, yes/no answers and other small talk

Two Ways to Do This

- Zero-Shot Text Classification

Zero-shot means:

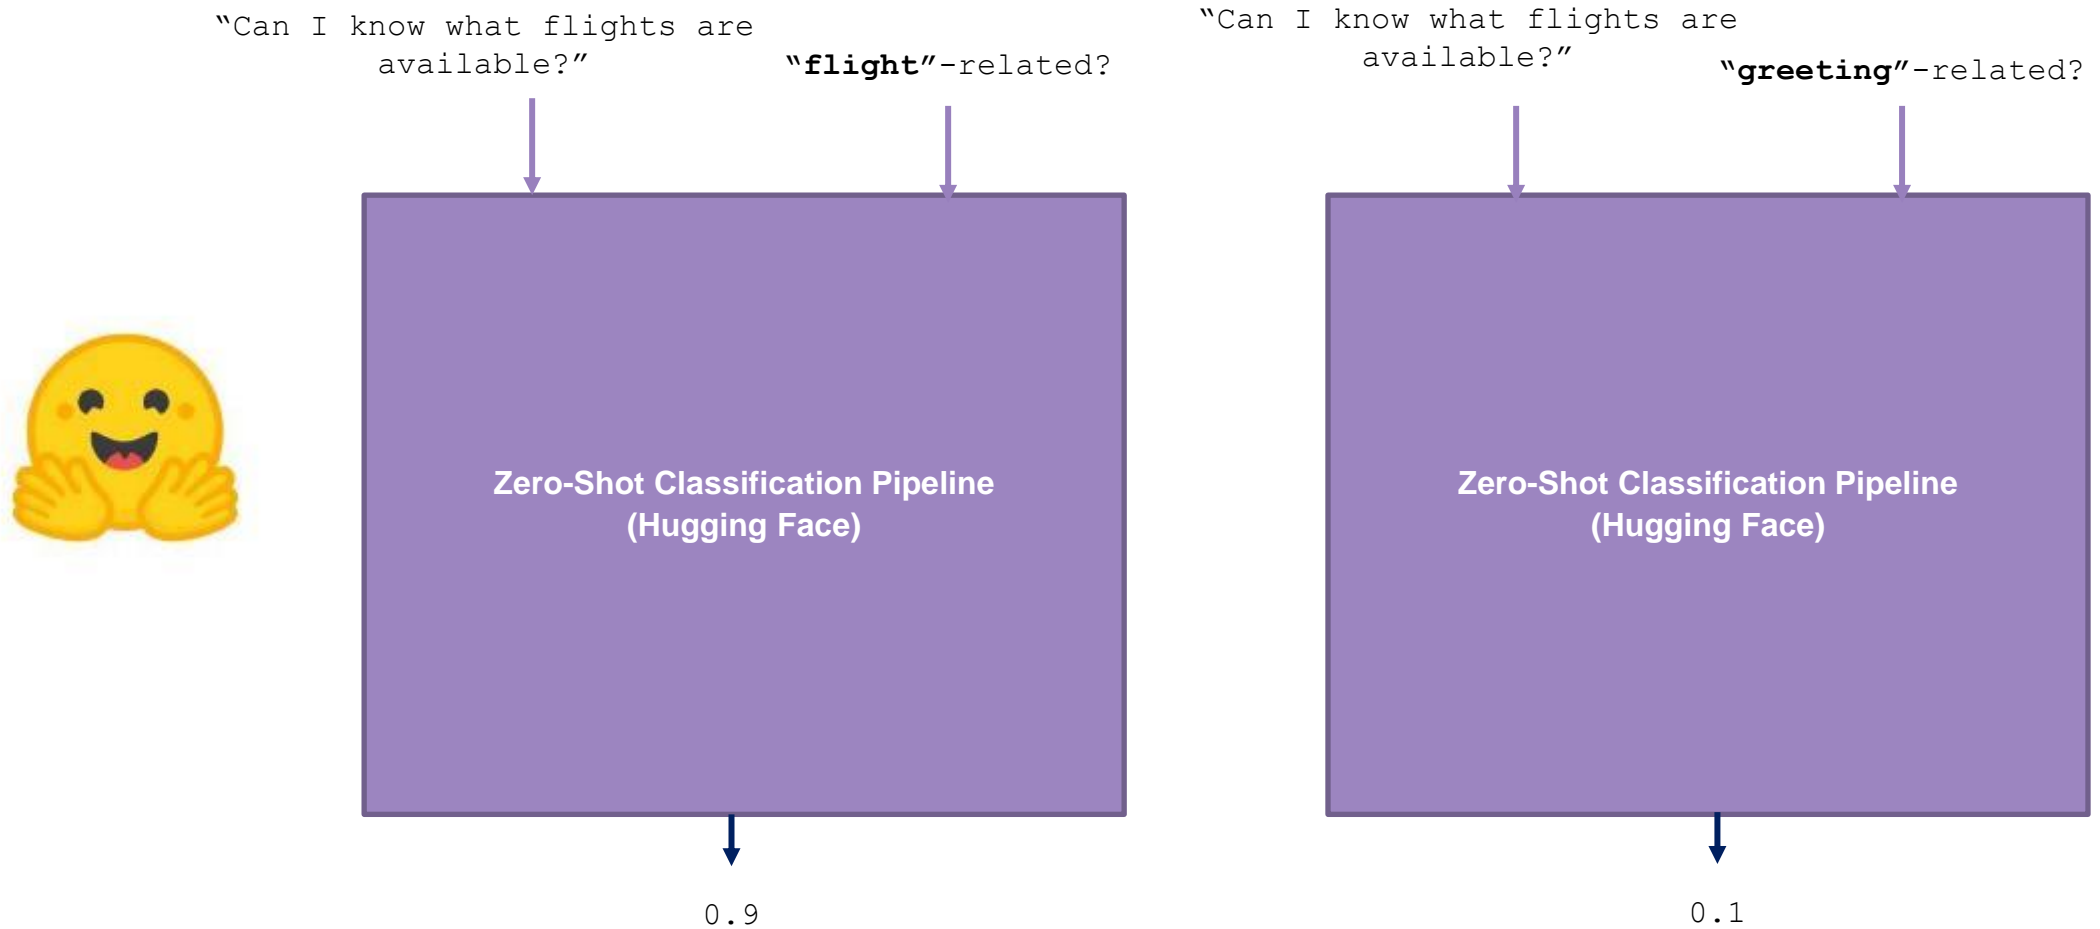
Use a pre-trained model with no *additional* training data for any generic task (this task was never seen in the original training data set).

This is good if you have no training data. Less accurate than training with your own custom dataset

<https://twitter.com/huggingface/status/1293240692924452864?lang=en>

- Collect Data and Perform Transfer Learning on the Pre-Trained Model with your Custom Labels


Transformers with HuggingFace – Zero-Shot



Highest score wins

Transformers with HuggingFace – Zero-Shot

Let's Try!!



Update: Zero-shot classification is now supported in our API and you can experiment with a number of compatible models on our [Model Hub](#).

Recently, the NLP science community has begun to pay increasing attention to zero-shot and few-shot applications, such as in the [paper from OpenAI](#) introducing GPT-3. This demo shows how 🤖 Transformers can be used for zero-shot topic classification, the task of predicting a topic that the model has not been trained on.

Zero Shot Topic Classification

Choose an example

"Jupyter's Biggest Moons Started as Tiny Grains of Hail"

Text

Jupiter's Biggest Moons Started as Tiny Grains of Hail

A new model offers an explanation for how the Galilean satellites formed around the solar system's largest world.

Konstantin Batygin did not set out to solve one of the solar system's most puzzling mysteries when he went for a run up a hill in Nice, France. Dr. Batygin, a Caltech researcher, best

Possible topics (separated by ',')

Go ahead and try different sets of broad, logically separate topics:

1. science, mathematics, English
2. politics, sports, finance
3. artificial intelligence, biology
4. happy, sad, neutral, angry

<https://huggingface.co/zero-shot/>

With Zero-Shot Learning

- The candidate “label” are English labels, and the model tries to compute the similarity the input text’s meaning with the label’s meaning.
- So to achieve good results, the choice of the candidate labels (with clear semantic boundary with the rest of the labels) is very important.

Creating Your Own Chatbot

INTENT RECOGNITION – TRAINING WITH CUSTOM LABELS

Transformers with HuggingFace

Custom Trained Labels

"Can I know what flights are available?"



Transformer Model
(Hugging Face)



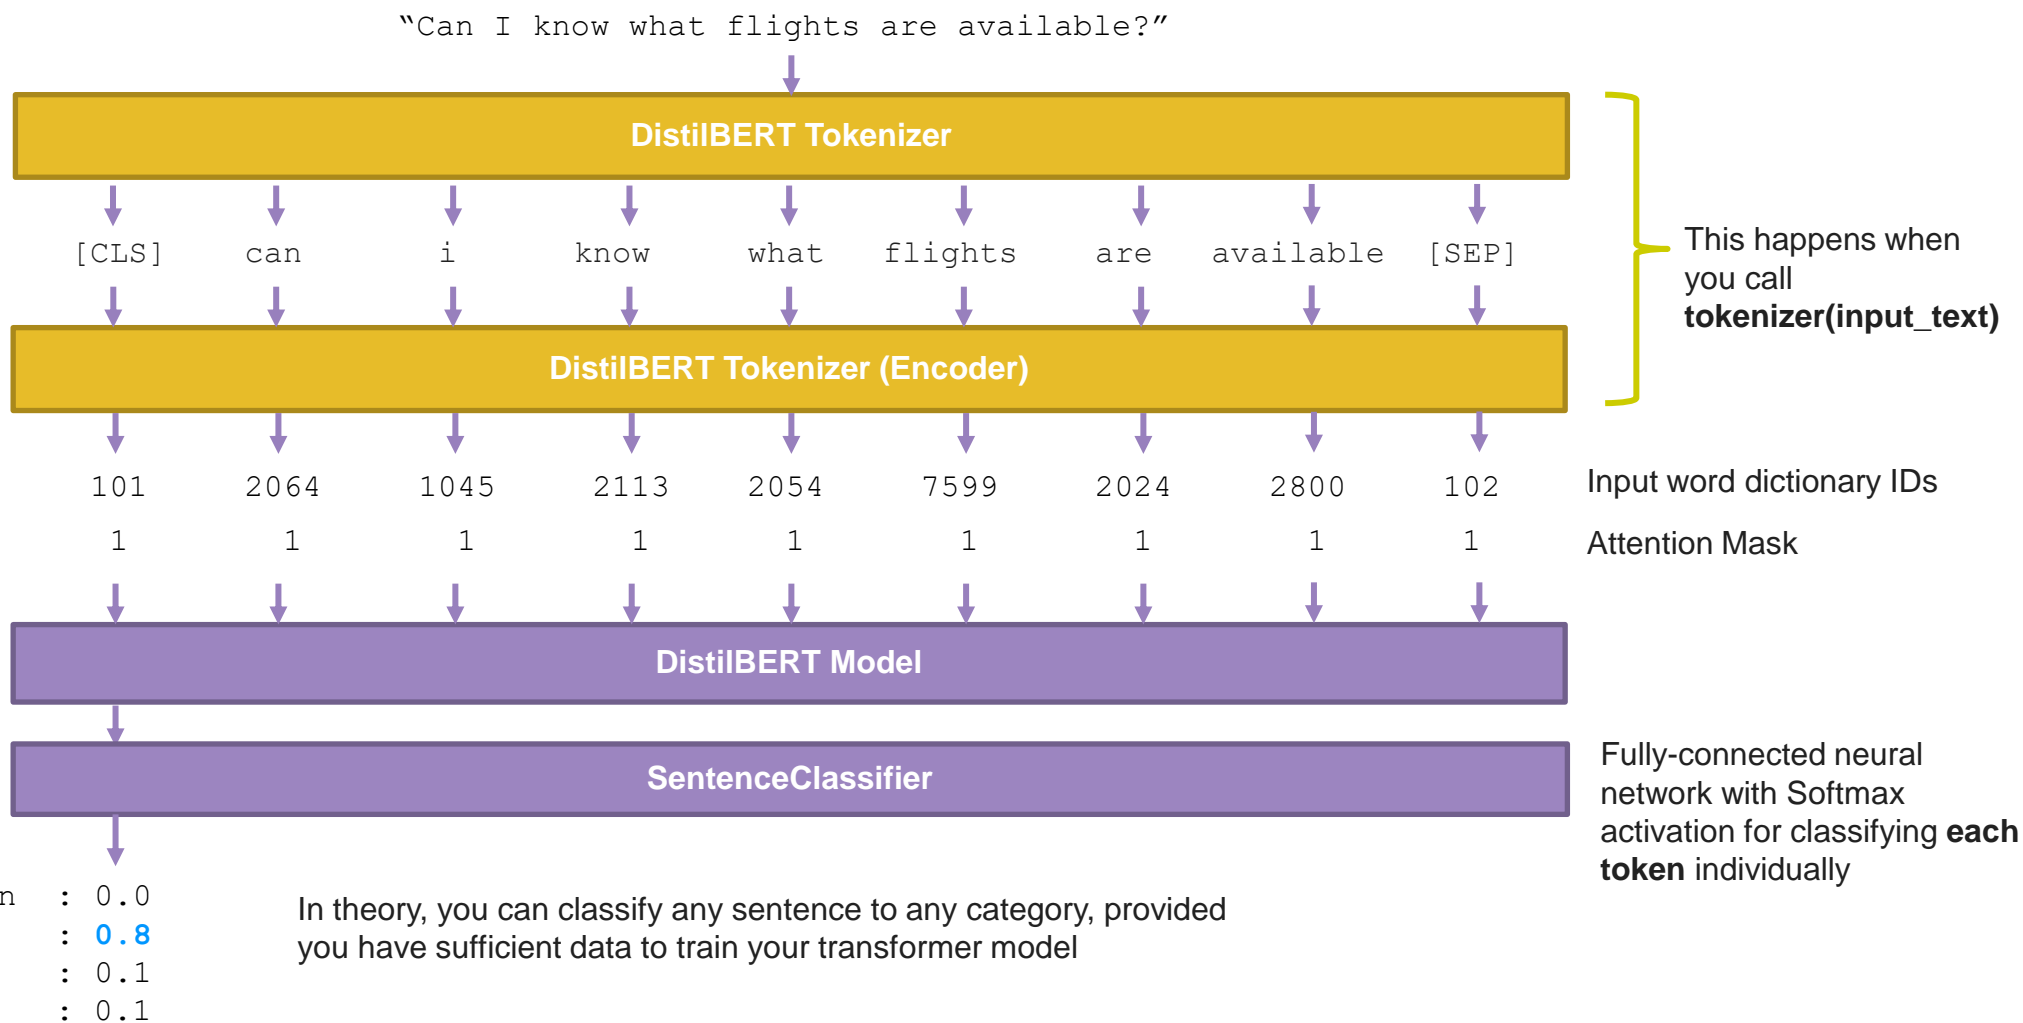
```
chat_abbreviation : 0.0
chat_flight       : 0.8
chat_yes          : 0.1
others            : 0.1
```

This is what we want to achieve:

- Pass in a sentence as input,
- Retrieve a classification as output

Transformers with HuggingFace

Custom Trained Labels

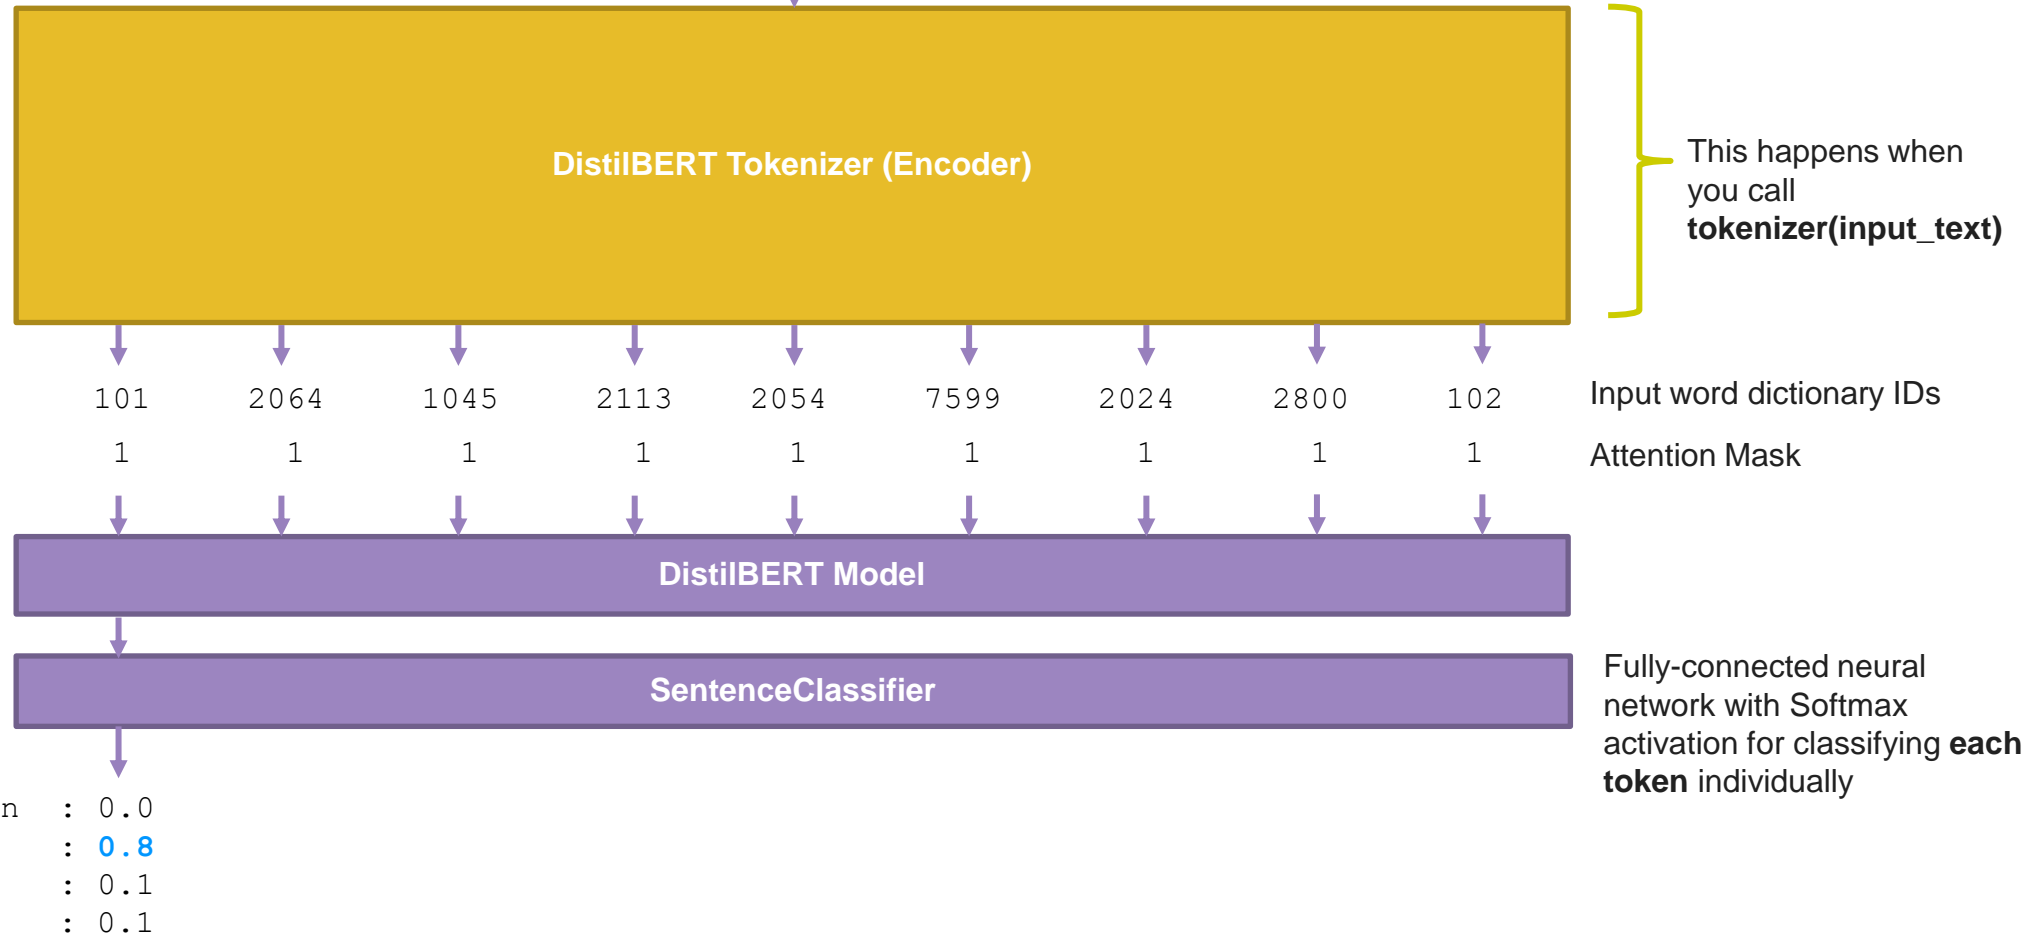


Transformers with HuggingFace

Custom Trained Labels



"Can I know what flights are available?"



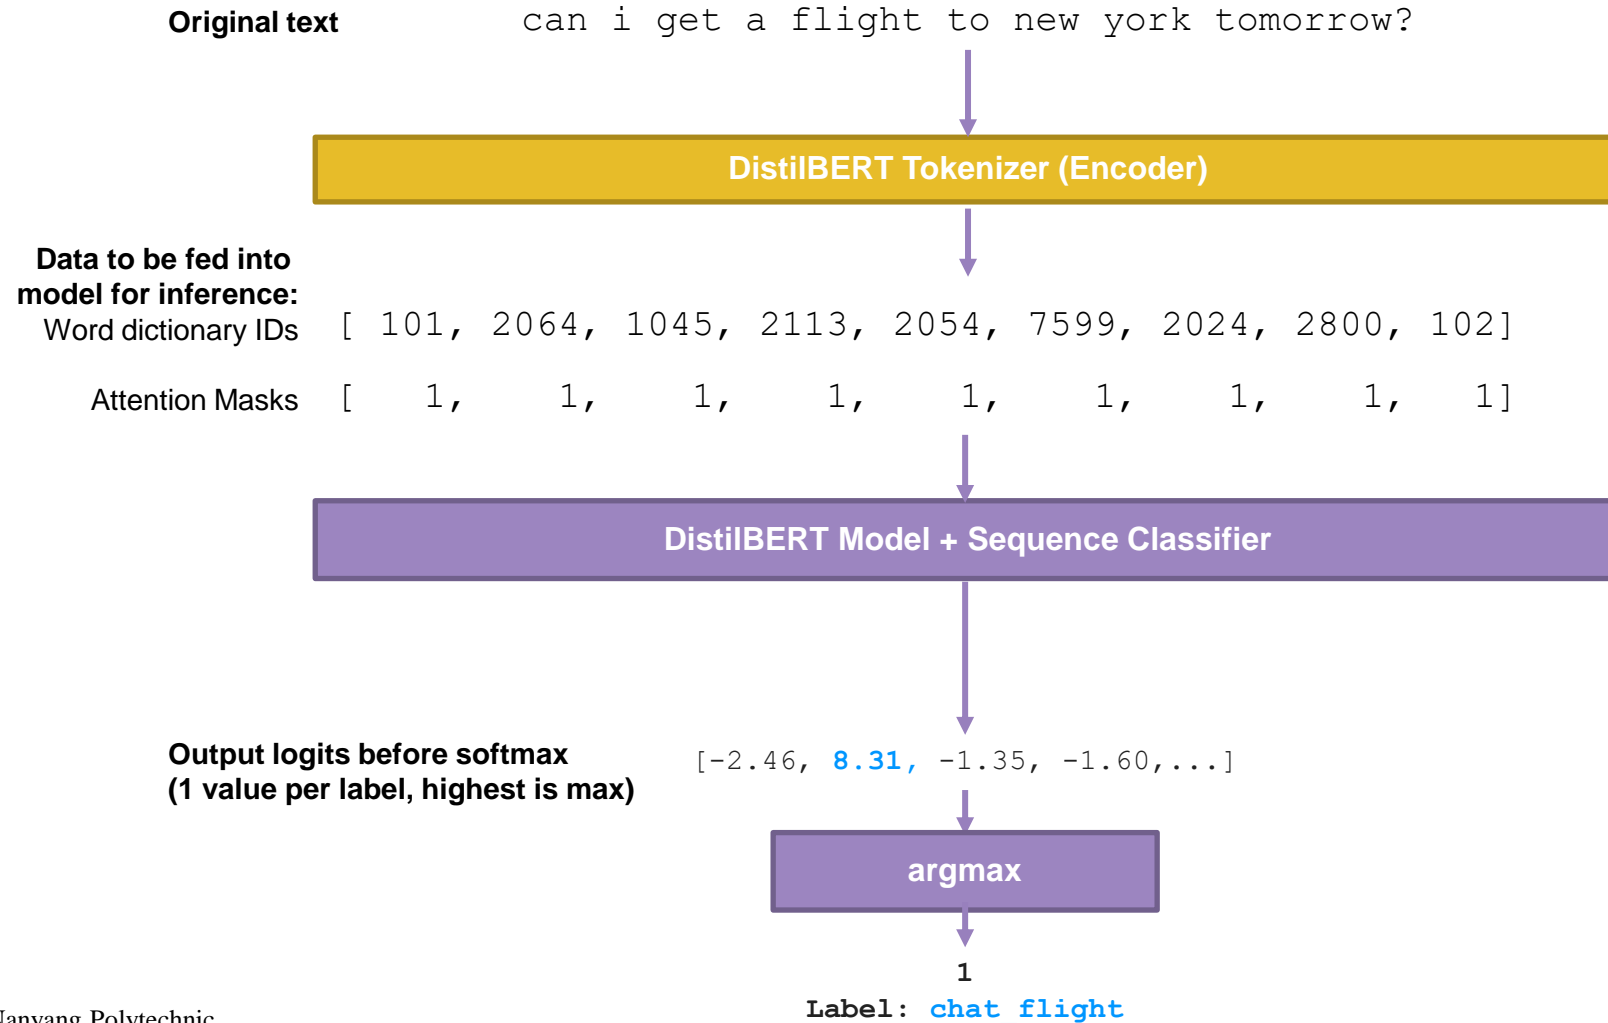
Training and Test Dataset

- We are going to use all of the `atis_intent.csv` data, but only part of the `smalltalkkb.tsv` dataset.
- We will construct the dataset manually with Excel to hold the following:
 - All of `atis_intent.csv`
 - Only greetings / yes / no from `smalltalkkb.tsv`

Training and Test Data Format

- The format required for the training and evaluation:
- X (inputs):
 - Output from the tokenizer (input word IDs, and attention mask)
- Y (output):
 - Index to the label
(training: *not* one-hot encoding, inference: one-hot encoding)

Inference



Creating Your Own Chatbot

CUSTOM NAMED ENTITY RECOGNITION

What is Named Entity Recognition?

Seeks to locate and classify named entity mentions in unstructured text into pre-defined categories such as:

the person names, organizations, locations, medical codes, time expressions, quantities, monetary values, percentages, etc.

https://en.wikipedia.org/wiki/Named-entity_recognition

What is Named Entity Recognition?

Jim bought 300 shares of Acme Corp. in 2006.

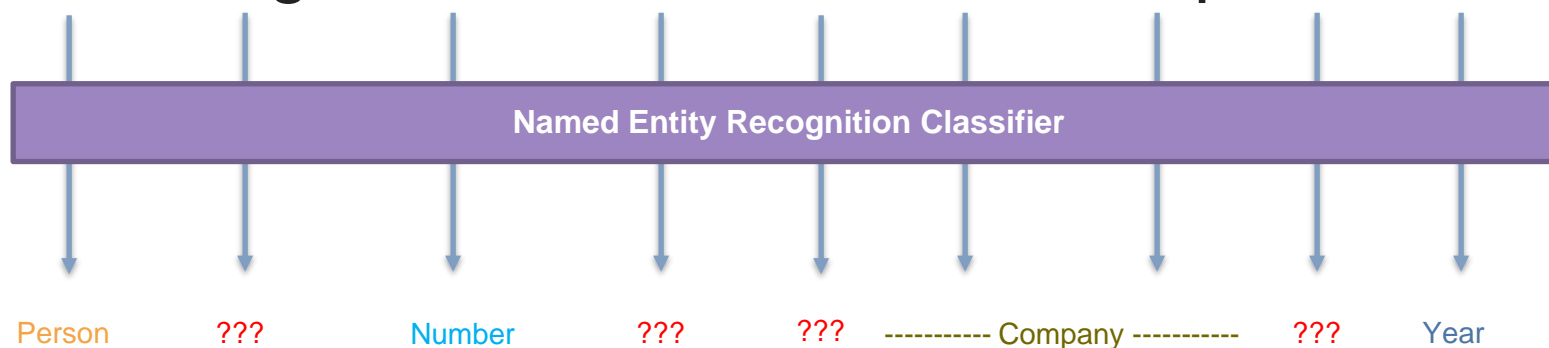
Person Number ----- Company ----- Year

DEMO

<https://cloud.google.com/natural-language>

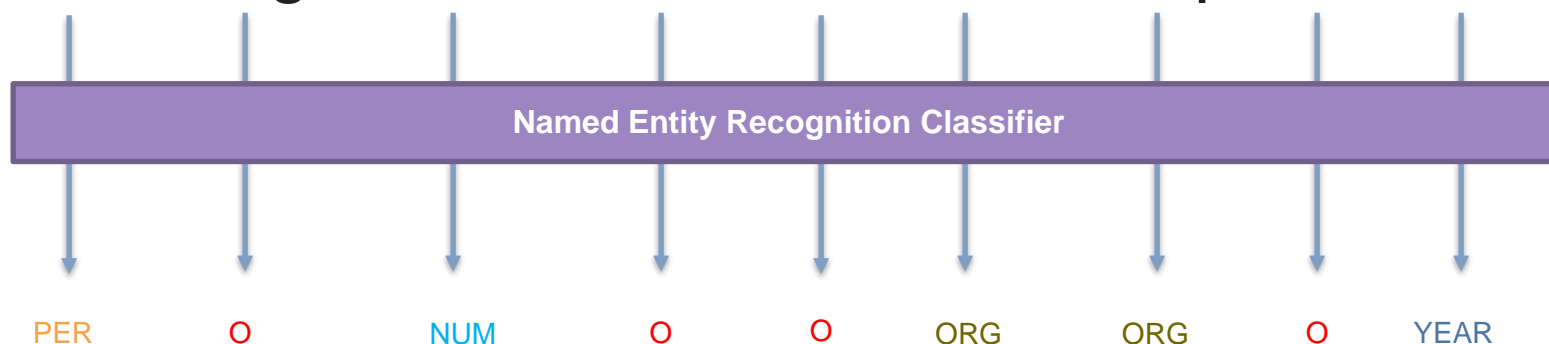
What is Named Entity Recognition?

Jim bought 300 shares of Acme Corp. in 2006.



What is Named Entity Recognition?

Jim bought 300 shares of Acme Corp. in 2006.

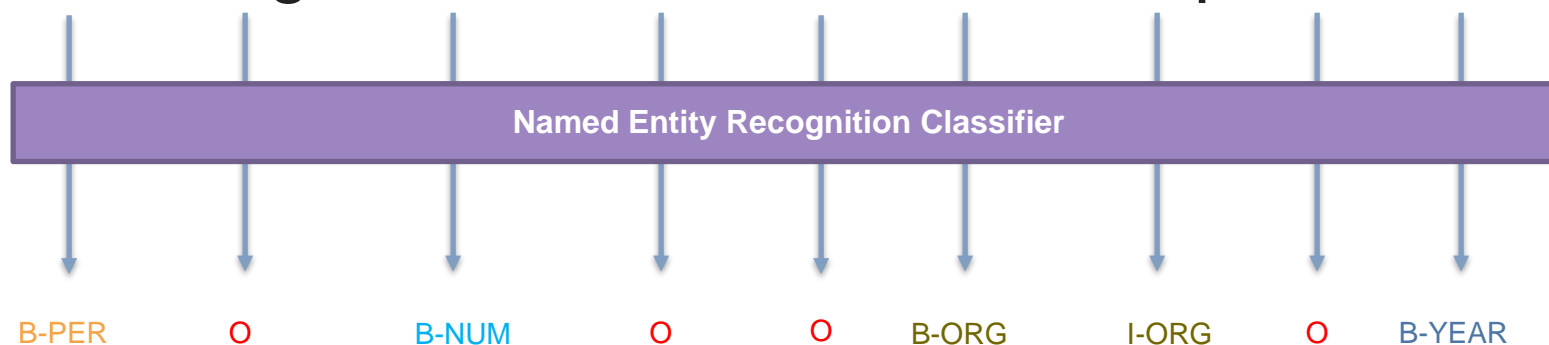


By convention, we give a code name for each entity we want to recognize.

“O” is a special entity that we will want to ignore.

NER – Used for Extracting Actionable Data

Jim bought 300 shares of Acme Corp. in 2006.



For multi-word entities like “Acme Corp”, we use:

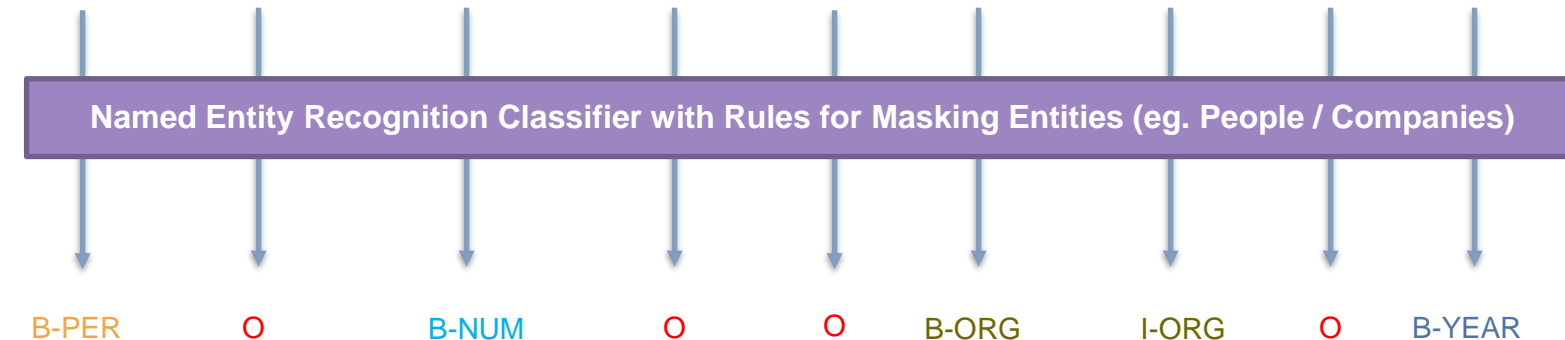
B-ORG (1st word), and

I-ORG (subsequent words joined to the first name)

This helps to extract important information that your app can act upon.

NER – Security and Personal Data Protection

Jim bought 300 shares of Acme Corp. in 2006.



***** bought 300 shares of ***** ***** in 2006.

Your app can also use the classified labels for each word to selectively mask them.

Use Cases – Information Extraction



Masking sensitive data (names, gender, age) from medical history



Masking sensitive data (names, banking information) from emails / online chats



Extracting important information from medical history



Extracting actionable information from users in a chatbot conversation

<https://youtu.be/6vX3Us1TOw8?t=159>


Difference with Entity Recognition?

- Remember Entity Recognition in DialogFlow?
- Named Entity Recognition is a pre-defined list of entities learned by the AI model.
- DialogFlow's Entity Recognition allow you to define custom entity types. For example,
I want to fix an appointment for repairing bicycle wheels



Custom entity with other possible values:
Handlebar, speedometer, bike frame, etc

Transformers with HuggingFace Again



HuggingFace... Again?

- Model creation:
 - Since this is used for classifying individual tokens (words) to try to recognize whether they are names of places: we must create a new, separate model.

The I- O- B- Format

- As mentioned, to label our entities, we use IOB format.
- Also known as **Inside-Outside-Beginning**:

Jim
bought
2
tickets
to
New
York

B-PER

O

O

O

O

B-LOC

I-LOC

This single word "Jim" forms the name of a person

The two words, "New York", form an instance of the location.

The first word of an entity is marked with B-###, second word onwards marked with I-###

Training and Test Dataset

- This is an example of the CoNLL 2003 dataset:
 - <https://github.com/synalp/NER/blob/master/corpus/CoNLL-2003/eng.testa>

-DOCSTART- -X- 0 0

CRICKET NNP I-NP 0

- : 0 0

LEICESTERSHIRE NNP I-NP I-ORG

TAKE NNP I-NP 0

OVER IN I-PP 0

AT NNP I-NP 0

TOP NNP I-NP 0

AFTER NNP I-NP 0

INNINGS NNP I-NP 0

VICTORY NN I-NP 0

. . 0 0

LONDON NNP I-NP I-LOC

1996-08-30 CD I-NP 0

5 Entity Types:

- | | |
|-------------------|------------------------------|
| - B-PER / I-PER | Person Name |
| - B-LOC / I-LOC | Country or Location Name |
| - B-ORG / I-ORG | Organization or Company Name |
| - B-MISC / I-MISC | Any Other Names |
| - O | Other Words to be Ignored |

Training and Test Data Format

- The format required for the training and evaluation:
- X (inputs):
 - Output from the tokenizer (input word IDs, and attention mask)
- Y (output):
 - A list of indexes to the individual word classification
(training: *not* one-hot encoding, inference: one-hot encoding)

Inference

Original text

need a flight to new york tomorrow

[CLS]

[SEP]

DistilBERT Tokenizer (Encoder)

Input IDs [101, 2064, 1045, 2113, 2054, 7599, 2024, 2800, 102]

Data to be fed into model for inference

Attention Mask [1, 1, 1, 1, 1, 1, 1, 1, 1]

DistilBERT Model + Token Classifier

O	:	1.0	1.0	1.0	1.0	1.0	0.1	0.0	1.0	1.0
B-LOC	:	0.0	0.0	0.0	0.0	0.0	0.9	0.0	0.0	0.0
I-LOC	:	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0

Output from the token classifier (softmax)

argmax

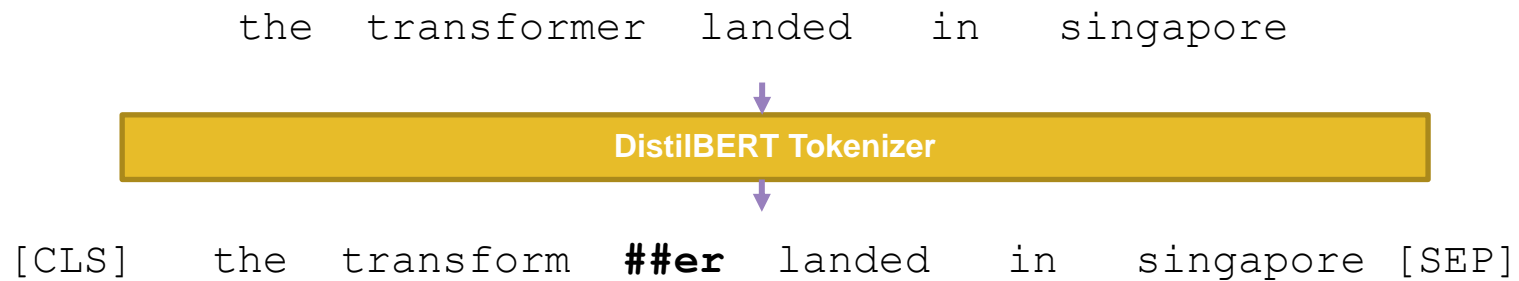
Label:	0	0	0	0	0	1	2	0	0
	O	O	O	O	O	B-LOC	I-LOC	O	O

Handling Sub-words

- The DistilBERT Transformers is trained on a dictionary of about 30,000+ words
- Not all words can exist in the dictionary, some words are represented with multiple tokens of sub-words.
- For example, the word
 - “transformer” is tokenized to 2 words:
 - “transform” + “##er”

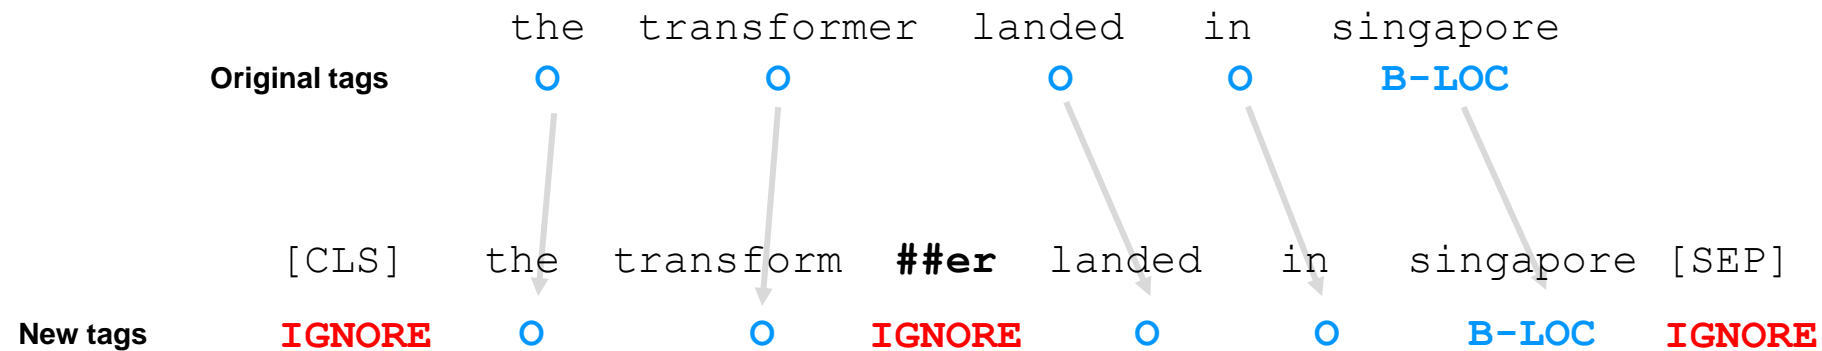
Handling Sub-words

- Big idea:
 - ignore the token labelling and classification for the sub-words.



Handling Sub-words

- Big idea:
 - ignore the token labelling and classification for the sub-words.



Inference

Original text

the transformer landed in singapore

Split into words

["the", "transformer", "landed", "in", "singapore"]

DistilBERT Tokenizer (Encoder with return_offset_mapping = True)

	[CLS]	the	transform	##er	landed	in	singapore	[SEP]
Input IDs	[101,	1996,	10938,	2121,	5565,	1999,	5264,	102]
Attention Mask	[1,	1,	1,	1,	1,	1,	1,	1]
Offset Mapping	[(0, 0),	(0, 3),	(0, 9),	(9, 11),	(0, 6),	(0, 2),	(0, 9),	(0, 0)]

Only Input IDs + Attention Mask are fed into model for inference

DistilBERT Model + Token Classifier

O	:	1.0	1.0	1.0	1.0	1.0	0.0	1.0
B-LOC	:	0.0	0.0	0.0	0.0	0.0	1.0	0.0
I-LOC	:	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Output from the token classifier (softmax)

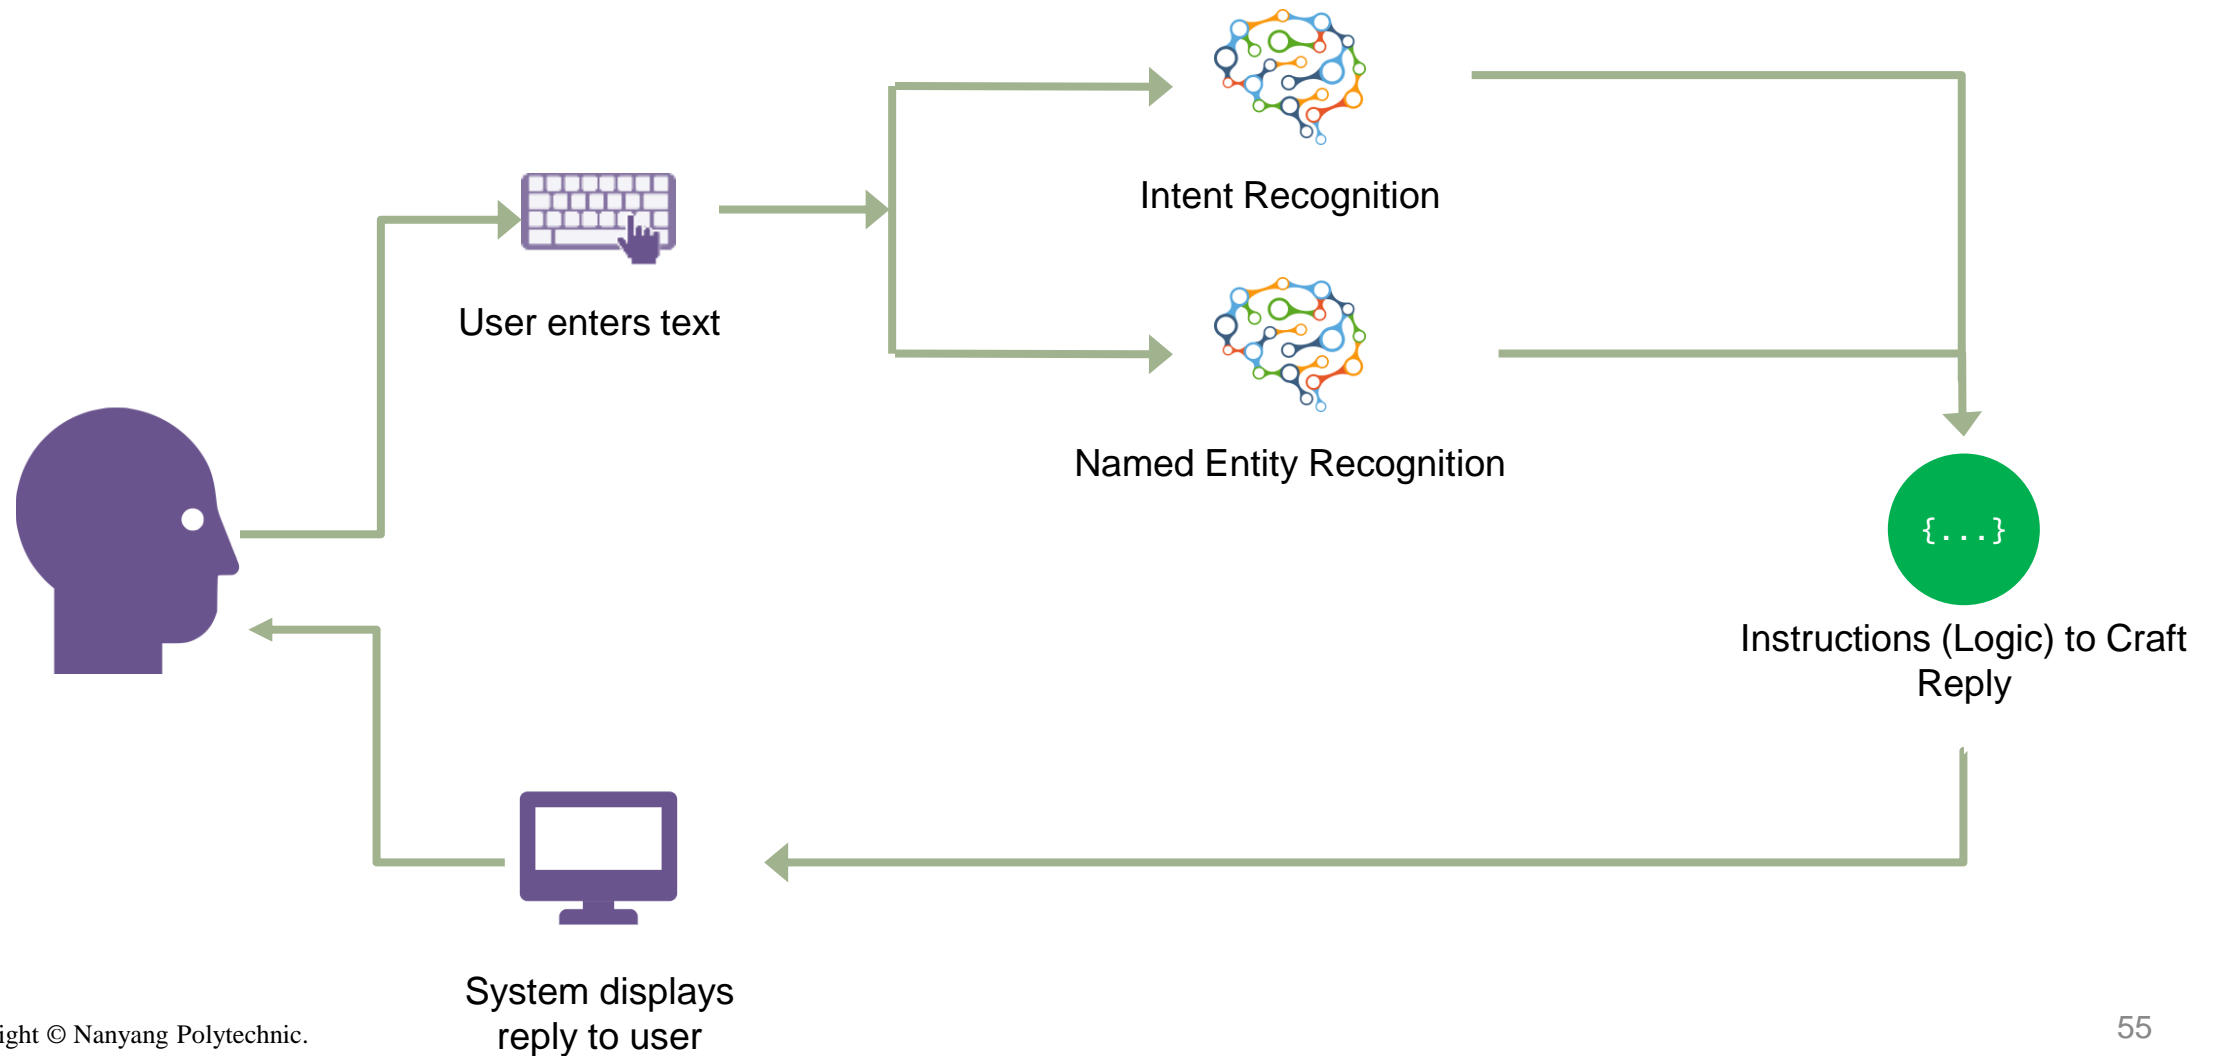
argmax

\emptyset 0 0 \emptyset 0 1 2 \emptyset

Creating Your Own Chatbot

PUTTING IT ALTOGETHER

Putting It Altogether



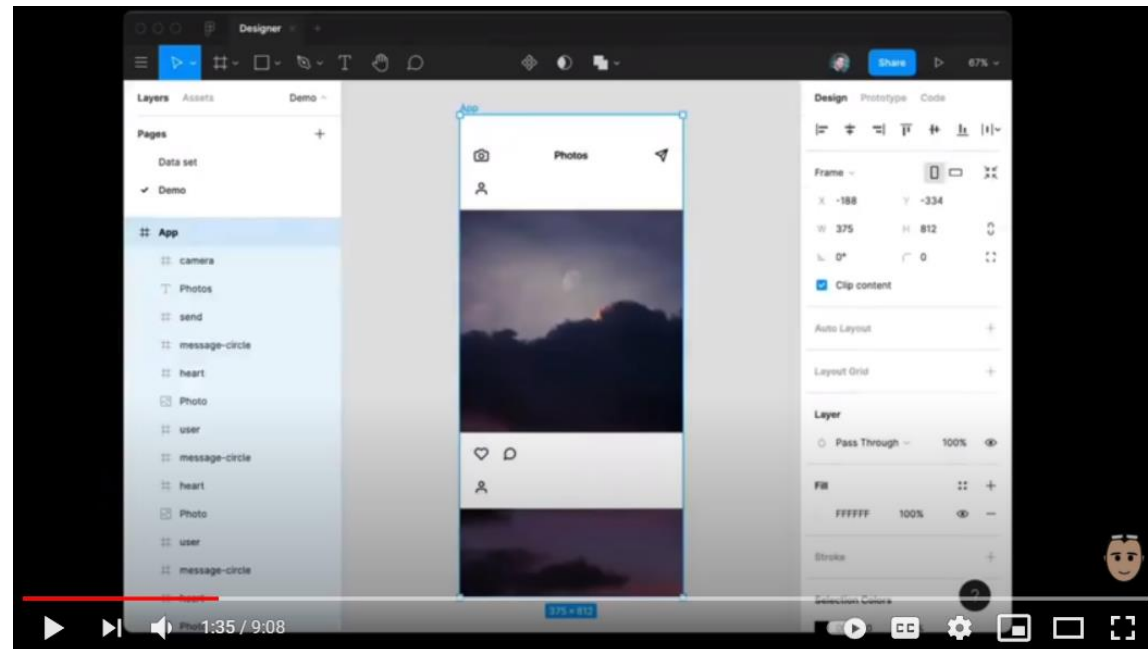
Creating Your Own Chatbot

OTHER ADVANCES IN NATURAL LANGUAGE PROCESSING

Other Advances in Natural Language Processing

OpenAI's GPT-3

<https://en.wikipedia.org/wiki/GPT-3>




<https://www.youtube.com/watch?v=8psgEDhT1MM>

Other Advances in Natural Language Processing

DALL.E (based on GPT-3)

<https://daleonai.com/dalle-5-mins>



<https://www.youtube.com/watch?v=8psgEDhT1MM>

Summary

- **What is a Chatbot?**
- **Transformers with the HuggingFace Library**
- **Intent Recognition**
- **Custom Named Entity Recognition**
- **Putting It Altogether**
- **Other Advances in Natural Language Processing**